

Class8

David Majeed (A17885958)

Table of contents

Intro	1
Data Import	1
Exploring the Data	3
PCA	4
Variance	7
Hierarchical Clustering	9
Combining methods	12
Sensitivity/ Specificity	14
Prediction	14

Intro

Our goal is to analysis data coming from breast cancer research using the techniques we have learnt in other classes.

Data Import

The data comes from fine needle aspiration

```
library(readr)

fna.data <- read_csv("Class8/WisconsinCancer.csv")
```

```
Rows: 569 Columns: 32
-- Column specification -----
Delimiter: ","
```

```
chr (1): diagnosis
dbl (31): id, radius_mean, texture_mean, perimeter_mean, area_mean, smoothne...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
#read.csv allows us to import our data set
```

```
wisc.df<- read.csv("Class8/WisconsinCancer.csv", row.names=1)
#Shift the row names by 1 to line up our data properly
```

```
head(wisc.df,3)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.8	1001
842517	M	20.57	17.77	132.9	1326
84300903	M	19.69	21.25	130.0	1203
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001		0.14710
842517	0.08474	0.07864	0.0869		0.07017
84300903	0.10960	0.15990	0.1974		0.12790
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003		0.006193	25.38	17.33
842517	0.01389		0.003532	24.99	23.41
84300903	0.02250		0.004571	23.57	25.53
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.6	2019	0.1622		0.6656
842517	158.8	1956	0.1238		0.1866
84300903	152.5	1709	0.1444		0.4245
	concavity_worst	concave.points_worst	symmetry_worst		
842302	0.7119		0.2654	0.4601	
842517	0.2416		0.1860	0.2750	
84300903	0.4504		0.2430	0.3613	
	fractal_dimension_worst				

842302	0.11890
842517	0.08902
84300903	0.08758

```
#Not to show up all the data

#lets remove the diagonsis column
wisc.data<- wisc.df[,-1]
#make it's own
diagonsis<- wisc.df[,1]
```

Exploring the Data

Q1. There are 569 observation

```
nrow(wisc.data)
```

```
[1] 569
```

Q2. There are 212 malignant diagnoses

```
#my attempt
sum(diagonsis=='M')
```

```
[1] 212
```

```
#Prof's way
table(diagonsis)
```

```
diagonsis
  B   M
357 212
```

Q3. There are 10 variables in the data with the suffix of __mean

```
length(grep("__mean", colnames(wisc.data), value=F))
```

```
[1] 10
```

```
#length counts
#grep searches for the suffix
```

PCA

Q4. The original variance is 0.4427 from PC1

Q5. PC3 is required to describe at least 70% of the original variance

Q6. PC7 is needed for at least 90% of the original variance data


```
#we need to scale our data
wisc.pr<-prcomp(wisc.data, scale=T)
summary(wisc.pr)
```

Importance of components:


	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q7. The thing that stands out is the lack of readability, all the labels and data points are stacked one another thus I can not understand it at all

```
biplot(wisc.pr)
```




```
#lets make a PC score plot
library(ggplot2)
ggplot(wisc.pr$x)+
  aes(PC1, PC2, col=diagonsis)+
  geom_point()
```



Q8. I notice that it has shifted everything down a bit. It still looks pretty similar to the previous plot but with more mixing of the two types together.

```
ggplot(wisc.pr$x)+  
  aes(PC1, PC3, col=diagnosis)+  
  geom_point()
```




Variance

```
pr.var<- wisc.pr$sdev^2
head(pr.var)
```


```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
# Variance explained by each principal component: pve
pve <- pr.var/sum(pr.var)

# Plot variance explained for each principal component
plot(c(1,pve), xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```

Communicating PCA results

```
wisc.pr$rotation["concave.points_mean",1]
```

```
[1] -0.2608538
```

Q.9 The loading value of the first principal component is -0.2608538. There are no other features that contribute more than this one.

Hierarchical Clustering

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)


#Find the distance
data.dist <- dist(data.scaled)

#Make the hclust
wisc.hclust <- hclust(data.dist,method="complete")
```

```
#make the plot
```

```
plot(wisc.hclust)  
abline(h = 19, col="red", lty=2)
```

Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

```
#Lets see what height can cut the tree into 4  
wisc.hclust.clusters <- cutree(wisc.hclust, h = 24)  
  
table(wisc.hclust.clusters)
```

```
wisc.hclust.clusters  
  1  2  
567  2
```

```
table(wisc.hclust.clusters, fna.data$diagnosis)
```

```
wisc.hclust.clusters   B   M  
      1 357 210  
      2   0   2
```

Q.10 The height at which the clustering model has 4 branches is at 19

Q.11 The height for 6 clusters is 16.9 and for 2 clusters the height is 24. Either way the data in between is not the best as it doesn't help us with narrowing down the data either being too little or too much


```
wisc.hclust <- hclust(data.dist,method="ward.D2")  
#changing the method  
  
wisc.hclust
```

Call:

```
hclust(d = data.dist, method = "ward.D2")
```

```
Cluster method   : ward.D2  
Distance          : euclidean  
Number of objects: 569
```

```
plot(wisc.hclust)
```




Q.12 The “ward.D2” is the better method to display the data as it gives a separation of the data unlike other methods, this allows us to see the different branches easier. With methods

like “complete”, there was a need for more guess work to figure out where it was optimal to cut.

Combining methods


```
d<-dist(wisc.pr$x[,1:4])
wisc.pr.hclust<-hclust(d, method="ward.D2")
plot(wisc.pr.hclust)
abline(h=70, col="red")
```



```
grps<-cutree(wisc.pr.hclust, h=70)
table(grps)
```

```
grps
  1  2
171 398
```

```
ggplot(wisc.pr$x) +
  aes(PC1, PC2) +
  geom_point(col=grps)
```



```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters, fna.data$diagnosis)
```

```
wisc.pr.hclust.clusters   B    M
                        1    6 165
                        2 351   47
```

```
table(wisc.hclust.clusters,diagnosis)
```

```
              diagnosis
wisc.hclust.clusters  B    M
                    1 357 210
                    2   0   2
```

Q.13 It separates out the clusters pretty well, with many true positives and true negatives, with relatively few false negatives or false positives.

Q.14 The data is not as clean as the previous one, with more overlap and splitting of the data into four groups rather than the proper separation of the groups into 2 benign and malignant

How does this clustering pattern correspond to expert diagnosis

```
table(diagnosis, grps)
```

```
      grps
diagnosis 1  2
B      6 351
M     165  47
```

Sensitivity/ Specificity

```
#Sensitivity: TP/(TP+FN)
188/(188+24)
```

```
[1] 0.8867925
```

```
#Specificity: TN/(TN+FP)
329/(329+28)
```

```
[1] 0.9215686
```

Q.15 The PCA based hierarchical was the best with a sensitivity of 88.7% and a specificity of 92.2%

Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```


```
      PC1      PC2      PC3      PC4      PC5      PC6      PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
      PC8      PC9      PC10      PC11      PC12      PC13      PC14
[1,] -0.2307350  0.1029569 -0.9272861  0.3411457  0.375921  0.1610764  1.187882
[2,] -0.3307423  0.5281896 -0.4855301  0.7173233 -1.185917  0.5893856  0.303029
```

	PC15	PC16	PC17	PC18	PC19	PC20
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500

	PC21	PC22	PC23	PC24	PC25	PC26
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238

	PC27	PC28	PC29	PC30
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q.16 We should prioritize patient group #2, just like the our previous data, we saw that benign patients are closer to one another and have less spread in comparison to malignant which are more spread out

```
sessionInfo()
```

```
R version 4.5.3 (2026-03-11 ucrt)
Platform: x86_64-w64-mingw32/x64
```

Running under: Windows 11 x64 (build 26200)

Matrix products: default

LAPACK version 3.12.1

locale:

[1] LC_COLLATE=English_United States.utf8
[2] LC_CTYPE=English_United States.utf8
[3] LC_MONETARY=English_United States.utf8
[4] LC_NUMERIC=C
[5] LC_TIME=English_United States.utf8

time zone: America/Los_Angeles

tzcode source: internal

attached base packages:

[1] stats graphics grDevices utils datasets methods base

other attached packages:

[1] ggplot2_4.0.2 readr_2.2.0

loaded via a namespace (and not attached):

[1] crayon_1.5.3	vctrs_0.7.2	cli_3.6.6	knitr_1.51
[5] rlang_1.2.0	xfun_0.57	generics_0.1.4	S7_0.2.1-1
[9] jsonlite_2.0.0	labeling_0.4.3	glue_1.8.0	bit_4.6.0
[13] htmltools_0.5.9	hms_1.1.4	scales_1.4.0	rmarkdown_2.31
[17] grid_4.5.3	evaluate_1.0.5	tibble_3.3.1	tzdb_0.5.0
[21] fastmap_1.2.0	yaml_2.3.12	lifecycle_1.0.5	compiler_4.5.3
[25] dplyr_1.2.1	RColorBrewer_1.1-3	pkgconfig_2.0.3	farver_2.1.2
[29] digest_0.6.39	R6_2.6.1	tidyselect_1.2.1	parallel_4.5.3
[33] vroom_1.7.1	pillar_1.11.1	magrittr_2.0.5	withr_3.0.2
[37] gtable_0.3.6	tools_4.5.3	bit64_4.8.0	